

Fast k-Nearest Neighbor Classifier (beta)

Contents

Fast Nearest Neighbor Searching	1
The FastKNN Classifier	1
Find the Best k	1
Plot Classification Decision Boundary	1

Fast KNN with shrinkage estimator for the class membership probabilities

Fast Nearest Neighbor Searching

The `fastknn` method implements a k-Nearest Neighbor (KNN) classifier based on the `ANN` library. ANN is written in C++ and is able to find the k nearest neighbors for every point in a given dataset in $O(N \log N)$ time. The package `RANN` provides an easy interface to use ANN library in R.

The FastKNN Classifier

The `fastknn` was developed to deal with very large datasets ($> 100k$ rows) and is ideal to `Kaggle` competitions. It can be up to 20x faster than the popular `knn` method from the R package `class`, for large datasets. Moreover, `fastknn` provides a shrinkage estimator to the class membership probabilities, based on the inverse distances of the nearest neighbors (**see the PDF version**):

$$P(x_i \in y_j) = \frac{\sum_{k=1}^K \left(\frac{1}{d_{ik}} \cdot (n_{ik} \in y_j) \right)}{\sum_{k=1}^K \left(\frac{1}{d_{ik}} \right)}$$

where x_i is the i^{th} test instance, y_j is the j^{th} unique class label, n_{ik} is the k^{th} nearest neighbor of x_i , and d_{ik} is the distance between x_i and n_{ik} . This estimator can be thought of as a weighted voting rule, where those neighbors that are more close to x_i will have more influence on predicting x_i 's label.

In general, the weighted estimator provides more **calibrated probabilities** when compared with the traditional estimator based on the label proportions of the nearest neighbors, and reduces **logarithmic loss** (log-loss).

How to use

Required Packages

Find the Best k

Plot Classification Decision Boundary